

Pillar 2: Spatial GraphRAG & The "Deja Vu" Seed System

Lightweight Hierarchical Topological Graphs for AI Navigation in Dynamic 3D Worlds

A pure-Python prototype demonstrating that AI agents can navigate massive, dynamic 3D worlds using seed-deterministic hierarchical graphs — achieving sub-millisecond queries, under 12 MB memory, and fewer than 20 LLM tokens.

Author: Pulsate Labs

Date: June 2026

Status: Prototype (Not Final)

Principal Investigator: Mohato Sefatsa

Table of Contents

1. Introduction	3
2. System Architecture	3
2.1 The Deja Vu Seed Principle	3
2.2 Two-Level Hierarchical Graph	4
2.3 Sparse-by-Design Schema	4
3. Methodology	4
3.1 Three-Phase Benchmark Protocol	4
3.2 Dual-Algorithm Pathfinding	5
3.3 Measurement Methodology	5
3.4 Four Hard Metrics	5
4. Results	6
4.1 All Hard Targets Passed	6
4.2 Headroom Analysis	6
4.3 Latency Performance	7
4.4 Memory Footprint	8
4.5 Token Payload Efficiency	8
4.6 Broken Bridge Mutation	9
4.7 Phase B: Local Discovery Lifecycle	9
5. Comparison with Other Navigation Systems	10
6. Test Environment	11
7. Plan Going Forward	12
7.1 Scale Testing	12
7.2 C++/Rust Production Implementation	12
7.3 LLM Integration	12
7.4 Dynamic World Scenarios	12
7.5 Multi-Agent Coordination	13

1. Introduction

Hypothesis: A lightweight, hierarchical topological graph seeded deterministically from a world seed can enable AI agents to navigate massive, dynamic 3D environments with sub-millisecond path queries, minimal memory footprints, and token payloads small enough to fit within any modern LLM context window, without relying on vision models, voxel matrices, or dense occupancy grids.

Modern game worlds and simulation environments span millions of cubic meters of navigable space. Traditional approaches to AI navigation rely on either monolithic NavMesh generation (computationally expensive to update), voxel-based representations (memory-hungry at scale), or vision-model planners that require expensive GPU inference for every navigation decision. Each of these paradigms struggles in at least one dimension: NavMeshes are brittle under dynamic world changes; voxel octrees consume tens to hundreds of megabytes even at moderate resolutions; and vision models incur high latency and token costs that make real-time path correction impractical.

Spatial GraphRAG proposes a fundamentally different architecture. Instead of storing dense spatial data or relying on perceptual inference, we construct a sparse, hierarchical topological graph from a deterministic world seed. The global layer (L0) provides highway routing across the entire world, while local sub-graphs (L1) are lazily inflated only around the agent's current position, providing fine-grained maneuvering detail exactly where needed. This is the "Deja Vu" seed system: because the graph is deterministic from the seed, the LLM already "knows" the map topology, and we only need to communicate compact waypoint instructions rather than full spatial descriptions.

The key insight is that navigation intelligence does not require perception intelligence. An agent does not need to "see" the world to navigate it efficiently; it needs a compact, queryable representation of connectivity. By making this representation seed-deterministic, we achieve two powerful properties simultaneously: zero-cost map sharing (the seed is a single integer) and instant reproducibility (any agent with the seed reconstructs the same graph). This report documents a prototype implementation and benchmark of this system, testing whether it can meet four hard performance targets that would make it viable for real-time AI navigation in production environments.

2. System Architecture

2.1 The Deja Vu Seed Principle

The name "Deja Vu" captures the central mechanism: the AI agent experiences the world as if it has seen it before, because it has, through the seed. A world seed (a single 32-bit integer) deterministically generates the entire global navigation graph via a seeded pseudo-random number generator. Every agent that receives the same seed constructs an identical graph, enabling zero-communication map synchronization. This eliminates the need to transmit map data between agents, between client and server, or between sessions. The seed itself becomes the map.

In practice, the seed feeds into the graph construction process. L0 global nodes are placed in a 3D volume with coordinates generated from the seeded random number generator, ensuring identical placement across all instances. A spanning tree is constructed first to guarantee graph connectivity, then additional edges are added to create alternate routes and realistic topology. Because the entire process is a pure function of the seed, the graph is fully reproducible with no external state or data files required.

2.2 Two-Level Hierarchical Graph

The graph is organized into two distinct levels that serve different navigation needs. The L0 global highway layer consists of 100 high-level world nodes (in the default configuration) connected by a spanning tree plus additional alternate-route edges. This layer handles long-haul routing across the entire world. Edges are weighted by Euclidean distance in 3D space, incorporating terrain elevation differences. The L0 graph is always resident in memory and serves as the primary routing substrate.

The L1 local detail layer fans out from each L0 node with 20 local child nodes per cluster (in the default configuration). L1 nodes represent specific navigable points within the local area around their parent L0 node, such as building entrances, caves, resource veins, and other features. Intra-cluster edges connect nearby L1 nodes within the same cluster, providing local maneuvering paths. Critically, L1 sub-graphs are not part of global pathfinding; they are only inflated when the agent enters the corresponding L0 cell and pruned when it departs, following a lazy-load/eager-prune cycle that keeps memory bounded.

2.3 Sparse-by-Design Schema

Every byte in the graph representation must earn its place. The sparse-by-design principle mandates that nodes carry only the minimum attributes needed for navigation: a unique identifier, a human-readable label (composed of biome/feature names for interpretability), the hierarchy level (L0 or L1), and three floating-point coordinates (x, y, z). Edges carry only a weight (Euclidean distance, float) and a status string ("ACTIVE" or "DESTROYED"). No semantic tags, no visual properties, no occupancy data. This austerity is deliberate: it minimizes memory footprint, accelerates graph operations by reducing cache misses, and ensures that the token payload for LLM context injection remains minimal.

3. Methodology

3.1 Three-Phase Benchmark Protocol

The benchmark protocol is structured into three phases that progressively stress the system from initialization through normal operation to dynamic mutation. Phase A (Global Map Initialization) measures the one-time cost of constructing the full graph from the seed, including memory allocation at scale (10,000 nodes, 25,000 edges). It also measures pathfinding query latency on a 1,000-node graph and the global token payload. Phase B (Local Discovery) measures the lifecycle cost of inflating and pruning local sub-graphs, plus the local token payload with inflated clusters. Phase C (Broken Bridge) measures the worst-case cost of responding to a dynamic world change: an edge along the current path is severed (simulating a destroyed bridge, closed door, or newly impassable terrain), and the system must

recalculate the detour from scratch.

3.2 Dual-Algorithm Pathfinding

Both Dijkstra's algorithm and A* with a Euclidean heuristic are benchmarked in every phase. Dijkstra serves as the baseline: it guarantees shortest paths but explores nodes uniformly in all directions. A* augments this with a heuristic function $h(n)$ = Euclidean distance from node n to the destination, computed from the node's (x, y, z) coordinates. Because edge weights are also Euclidean distances, the heuristic is both admissible (never overestimates) and consistent (satisfies the triangle inequality), guaranteeing that A* finds optimal paths while exploring fewer nodes. The dual benchmark quantifies the practical benefit of the heuristic on a real graph topology.

3.3 Measurement Methodology

Memory is measured using Python's tracemalloc module, which tracks both current and peak allocated memory. The peak metric is reported because it captures the maximum memory pressure during graph construction, which is the relevant figure for capacity planning. Latency is measured using `time.perf_counter()` at sub-microsecond resolution. Each query is executed over 50 trials with randomized source/destination pairs, and the average latency is reported. Token payload is estimated using a conservative word-based approximation (~ 1.3 tokens per word for English), measuring the "Where am I?" context that would be injected into an LLM prompt.

3.4 Four Hard Metrics

The system is evaluated against four hard performance targets derived from practical constraints in game engines, robotics systems, and LLM context windows:

Metric	Target	Rationale
Memory (peak)	< 15 MB	Must fit alongside game engine assets; mobile devices have limited RAM
Query Latency	< 2 ms	Real-time navigation at 60 FPS allows < 16 ms per frame; path query must be a fraction of that budget
Mutation + Recalc	< 5 ms	Dynamic world changes (doors, bridges) require near-instant rerouting without frame hitches
Token Payload	< 50 tokens	LLM context windows are scarce; navigation instructions must not crowd out reasoning content

Table 1: Hard performance targets and their practical rationale.

4. Results

4.1 All Hard Targets Passed

The prototype simulator passed all four hard metrics with significant headroom. The full graph (10,000 nodes, 25,000 edges) was constructed from seed with a peak memory of 11.37 MB. Pathfinding queries across a 1,000-node graph complete in under 1.2 ms for both algorithms. Mutation and detour recalculation across a 500-node graph completes in under 0.7 ms. Token payloads for both global and local contexts remain well under 50 tokens. The complete results are summarized below:

Metric	Achieved	Target	Headroom	Verdict
Memory (peak)	11.37 MB	15 MB	+24%	PASS
Query Latency (Dijkstra)	1.1912 ms	2 ms	+40%	PASS
Query Latency (A*)	0.7189 ms	2 ms	+64%	PASS
Mutation + Recalc (Dijkstra)	0.6667 ms	5 ms	+87%	PASS
Mutation + Recalc (A*)	0.5229 ms	5 ms	+90%	PASS
Token Payload (global)	9.9	50	+80%	PASS
Token Payload (local)	19	50	+62%	PASS

Table 2: Complete benchmark results across all four hard metrics.

4.2 Headroom Analysis

The headroom chart below visualizes the percentage by which each metric beats its target. Memory shows 24% headroom, meaning the 10K-node graph uses roughly three-quarters of the allocated budget. Mutation recalculation shows the largest headroom at 87-89%, demonstrating that the system can handle dynamic world changes with an order of magnitude more speed than required. Token payloads show 62-80% headroom, confirming that navigation instructions consume only a small fraction of the available context window budget. Even the tightest metric, query latency for Dijkstra at 40% headroom, still delivers nearly twice the required speed.




Figure 1: Headroom below hard targets for all seven benchmark measurements.

4.3 Latency Performance

Path queries on the 1,000-node graph complete in 0.7189 ms (A*) to 1.1912 ms (Dijkstra), both well under the 2 ms target. A* with its Euclidean heuristic achieves a substantial speed advantage over Dijkstra, completing queries approximately 40% faster. This advantage stems from A* directing its search toward the destination using the Euclidean distance heuristic, pruning nodes that are geometrically far from the goal. On sparse graphs with long average edge lengths (as in our random geometric graphs), this pruning effect is particularly pronounced because the heuristic can eliminate large regions of the graph from consideration.




Figure 2: Query and mutation latency for both algorithms, with hard target lines shown.

The direct comparison between Dijkstra and A* reveals that A* is approximately 40% faster on our benchmark graph. Dijkstra averages 1.1912 ms per query while A* averages 0.7189 ms. Both are well within the 2 ms target, but the A* advantage becomes more significant at larger scales where Dijkstra's

exhaustive exploration would become prohibitive.




Figure 3: Dijkstra vs A* pathfinding latency comparison on the 1K-node graph.

4.4 Memory Footprint

The full graph (10,000 nodes, 25,000 edges) consumes 11.37 MB of peak memory, which is 24% below the 15 MB target. This is a direct consequence of the sparse-by-design schema: each node stores only a label string, a level string, and three float coordinates; each edge stores only a weight float and a status string. The networkx Graph object adds some internal overhead for adjacency dictionaries, but the total remains well within budget. For context, a voxel octree representing the same spatial extent at 1-meter resolution would require over 100 MB, making GraphRAG roughly 10x more memory-efficient.




Figure 4: Memory footprint comparison between achieved peak and hard target.

4.5 Token Payload Efficiency

The token payload represents the estimated number of LLM tokens required to inject navigation context into the agent's reasoning window. Two scenarios are measured: the global "Where am I?" payload (current node label plus immediate neighbor labels) averaging 9.9 tokens, and the local inflated payload (same node with its L1 cluster inflated) at 19 tokens. Both are far below the 50-token target, with 80% and 62% headroom respectively. The global payload is remarkably compact because the sparse graph means most nodes have only a handful of neighbors, and the label format ("You are at Biome_42. Connected: Plains_7, Forest_89, Mountain_3.") is inherently token-efficient.




Figure 5: Token payload for global and local navigation contexts vs target.

4.6 Broken Bridge Mutation

The Broken Bridge phase simulates the most demanding real-time scenario: the agent is following a known path when an edge is severed (representing a destroyed bridge, a closed door, or a newly impassable area), and the system must recalculate the optimal detour immediately. Our benchmark severs a random edge along the current path, removes it from the graph, and re-runs the full pathfinding algorithm. The results are striking: Dijkstra completes the sever+recalc in 0.6667 ms and A* in 0.5229 ms, both approximately 87-89% below the 5 ms target. This demonstrates that the system's mutation reflex is an order of magnitude faster than required, providing ample headroom for more complex mutation scenarios.

4.7 Phase B: Local Discovery Lifecycle

The inflate/prune lifecycle of local sub-graphs was benchmarked separately. Inflating a 20-node local cluster around a global node takes 0.62 ms and adds minimal memory (0.02 MB). Pruning the same cluster takes just 0.06 ms. These operations are fast enough to be performed every frame without impacting the navigation budget, confirming that the lazy-load/eager-prune design is practical for real-time use.




Figure 6: Phase B lifecycle latency for inflate and prune operations.

5. Comparison with Other Navigation Systems

To contextualize the Spatial GraphRAG results, we compare against four established navigation paradigms across the same metrics. The comparison values for alternative systems are drawn from published literature, open-source benchmarks, and industry experience; they represent typical rather than worst-case performance for each approach.

System	Memory	Query	Mutation	Tokens
Spatial GraphRAG (this work)	11.4 MB	0.72 ms	0.52 ms	10
NavMesh (Recast/Detour)	25 MB	0.5 ms	50-200 ms	85
Voxel Octree	120 MB	8 ms	5 ms	200
Vision Model (Nav)	500+ MB	50 ms	50 ms	450
Flat Grid A*	15 MB	3 ms	3 ms	120

Table 3: Comparison of navigation systems across key metrics.

NavMesh systems (such as Recast/Detour used in Unity and Unreal Engine) excel at query latency thanks to their pre-computed convex polygon regions, but they suffer during dynamic mutation: rebuilding a NavMesh tile after a world change can take 50-200 ms, which is two orders of magnitude slower than GraphRAG's instant detour recalculation. Voxel octrees provide fine-grained 3D occupancy data but at enormous memory cost (120 MB for a moderate world) and with query latencies that are significantly slower than GraphRAG due to tree traversal overhead. Vision model navigation, while flexible, requires GPU inference for every navigation decision, resulting in 50 ms latencies and 450+ token payloads, making it impractical for real-time path correction. Flat grid A* operates on the same

algorithmic foundation as GraphRAG but without hierarchy, forcing it to pathfind over the full grid, which increases both latency and token payload.




Figure 7: Normalized comparison of five navigation systems (lower = better).

The critical differentiator for Spatial GraphRAG is not any single metric but the combination: no other system simultaneously achieves low memory, fast queries, fast mutations, and small token payloads. NavMesh is fast but mutation-slow and token-heavy. Voxel octrees are thorough but resource-profligate. Vision models are flexible but latency-bound. Flat grids are simple but unscalable. Only the hierarchical, seed-deterministic, sparse-by-design approach of Spatial GraphRAG meets all four constraints simultaneously.

6. Test Environment

Parameter	Value
Language	Python 3.12
Graph Library	networkx (Graph construction + Dijkstra/A* pathfinding)
Memory Profiler	tracemalloc (current + peak allocation tracking)
Timer	time.perf_counter() (sub-microsecond resolution)
Scale Test	10,000 nodes / 25,000 edges
Query Test	1,000 nodes, 50 randomized trial pairs
Mutation Test	500 nodes, 50 trials (sever + recalc)
Token Estimation	~1.3 tokens/word (conservative English approximation)

World Seed	42 (deterministic)
------------	--------------------

Table 4: Test environment and benchmark configuration.

7. Plan Going Forward

The prototype has validated the core hypothesis and demonstrated that all four hard metrics are achievable with significant headroom. The next phase of development will focus on scaling, robustness, and integration:

7.1 Scale Testing

The current prototype tests up to 10,000 nodes. Production game worlds may require 100,000+ L0 nodes with rich local clusters. We plan to benchmark scaling behavior up to 100,000 nodes to verify that memory stays under 15 MB and query latency remains sub-2ms. The hierarchical design should scale favorably: query latency depends on the number of explored nodes, not the total graph size, and A* on a sparse graph with Euclidean heuristic explores a narrow corridor toward the destination. Memory scales linearly with node count, so 100,000 nodes would consume approximately 110 MB, requiring either graph partitioning or selective loading strategies.

7.2 C++/Rust Production Implementation

The Python prototype proves the concept but incurs interpreter overhead that would not exist in a compiled implementation. A C++ or Rust implementation would eliminate Python's object overhead (each dict entry, each float object, each string), reducing memory by an estimated 3-5x and cutting latency by 2-10x. Based on the Python prototype's sub-millisecond query latency, a Rust implementation targeting 100-500 microsecond queries is realistic, enabling path queries at 1000+ Hz for agents that need real-time obstacle avoidance.

7.3 LLM Integration

The token payload measurements use a conservative word-based estimation. The next step is to implement actual LLM integration, where the agent receives seed-derived graph context as part of its system prompt and uses it to make navigation decisions in a simulated environment. This will validate the assumption that 10-20 tokens of navigation context is sufficient for an LLM to reason about pathfinding, detour selection, and spatial planning without hallucinating non-existent routes.

7.4 Dynamic World Scenarios

The Broken Bridge benchmark tests a single edge sever. Real dynamic worlds experience multiple simultaneous changes: doors opening and closing, terrain deforming, moving obstacles, and player-built structures. We plan to test batch mutations (10-100 simultaneous edge changes), partial graph reconstruction (re-seeding specific L1 clusters), and time-varying edge weights (representing seasonal

terrain changes or traffic patterns). The mutation recalculation's 87-89% headroom suggests that even complex mutation scenarios will remain within the 5 ms target.

7.5 Multi-Agent Coordination

The seed-deterministic design naturally supports multi-agent scenarios: all agents sharing the same seed have identical graph representations without any communication overhead. When one agent discovers a world change (a severed edge), it can broadcast just the edge identifier to other agents, who can instantly update their local graph copy. This gossip-based synchronization model requires minimal bandwidth and scales to hundreds of agents. We plan to prototype this in a multi-agent simulation environment.

8. Prototype Disclaimer

IMPORTANT: This is a prototype, NOT the final system. The results presented in this report demonstrate the viability of the Spatial GraphRAG approach and validate the core hypothesis that lightweight hierarchical topological graphs can support real-time AI navigation. However, the current implementation has several known limitations that must be addressed before production deployment: it uses Python rather than a compiled language, which inflates both memory and latency; the graph topology uses random geometric graphs rather than domain-specific world layouts; the token payload estimation uses a simplified word-count model rather than actual LLM tokenizer measurement; and the mutation scenarios are limited to single-edge severing rather than complex dynamic world changes. The performance targets should be viewed as proof-of-concept milestones, not production guarantees. The final system will require rigorous engineering, optimization, and validation in a real game engine or simulation platform.